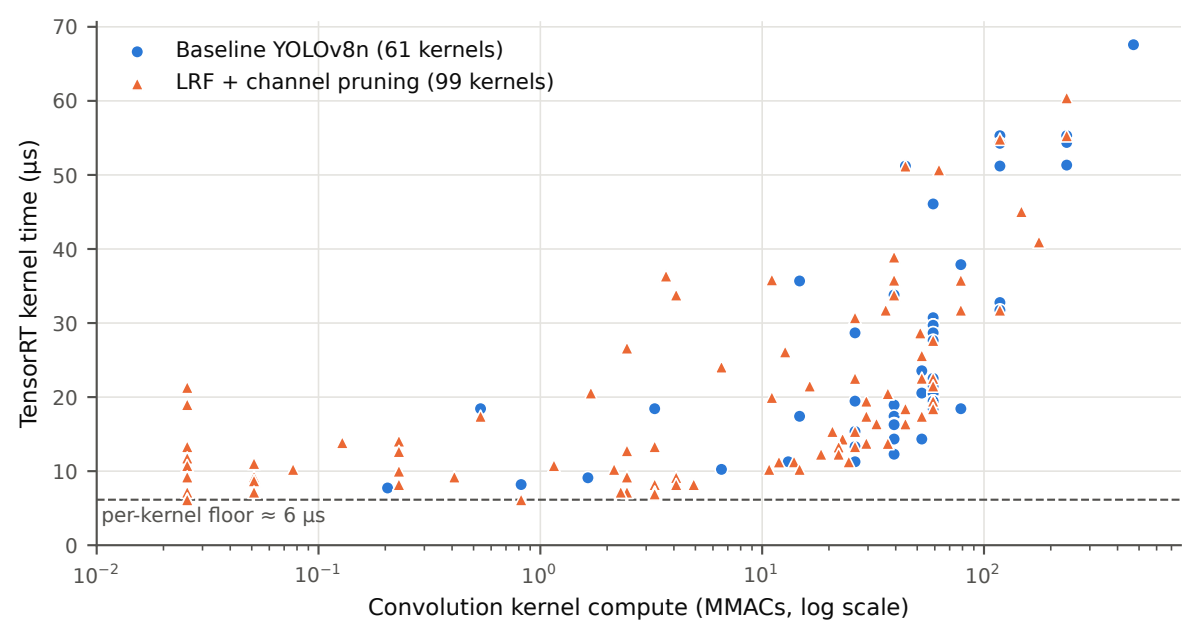
Baseline YOLOv8n (61 kernels)
LRF + channel pruning (99 kernels)
per-kernel floor ≈ 6 µs
TensorRT kernel time (µs)
Convolution kernel compute (MMACs, log scale)
0
10
20
30
40
50
60
70
10^{-2}
10^{-1}
10^{0}
10^{1}
10^{2}